

In [131]:

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
```

## a) load/merge data and visualize logerror

In [132]:

```
# load data into DataFrames
df1=pd.read_csv("train.csv")
df2=pd.read_csv("properties.csv")
df_final = df1.merge(df2, on = 'id', how = 'left')
```

In [133]:

```
# eliminate outliers
LE1=np.percentile(df_final.iloc[:,1],1,interpolation='nearest')
LE99=np.percentile(df_final.iloc[:,1],99,interpolation='nearest')
logerr=np.array(df_final.iloc[:,1])
logerr[np.where(logerr>=LE99)]=LE99
logerr[np.where(logerr<=LE1)]=LE1

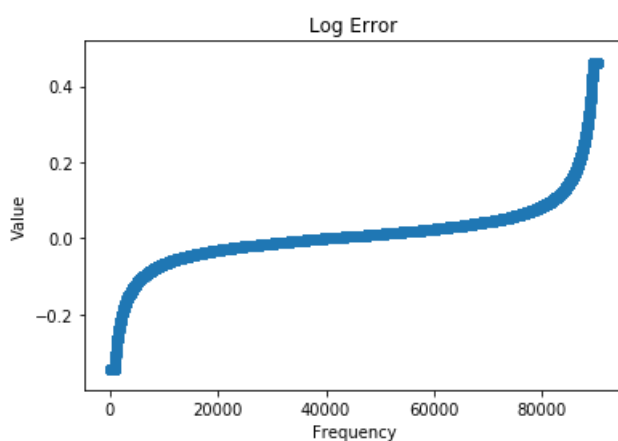
df_final.iloc[:,1]=logerr
```

In [134]:

```
# scatter of logerr
a = df_final.iloc[:,1]
plt.scatter(np.arange(len(df_final)), np.sort(a))
plt.xlabel('Frequency')
plt.ylabel('Value')
plt.title('Log Error')
```

Out[134]:

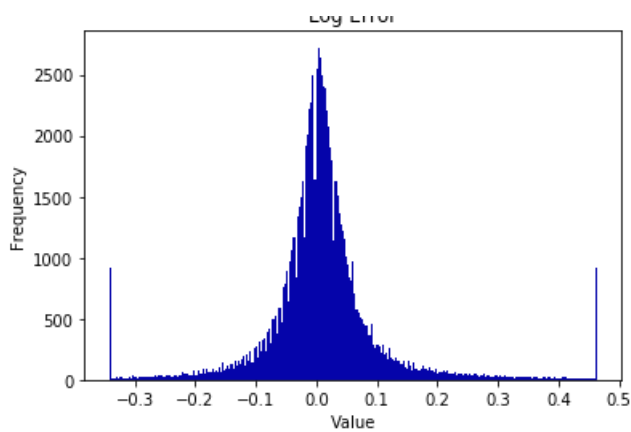
Text(0.5, 1.0, 'Log Error')



In [135]:

```
# histogram of logerr
n, bins, patches = plt.hist(df_final.iloc[:,1], bins='auto', color='#0504aa')
plt.xlabel('Value')
plt.ylabel('Frequency')
plt.title('Log Error')
maxfreq = n.max()
```

Log Error



## b) data cleaning

In [136]:

```
# build new data frame
df_1=np.array(df_final.columns)
df_2=np.array(df_final.isnull().sum(axis = 0))
df_11=np.column_stack((df_1,df_2))
df_12=pd.DataFrame(df_11)
df_12.columns = ["column_name", "missing_count"]
df_12['missing_ratio'] = df_12.iloc[:,1]/len(df_final)
print(df_12)
```

	column_name	missing_count	missing_ratio
0	id	0	0
1	logerror	0	0
2	transactiondate	0	0
3	airconditioningtypeid	80113	0.887433
4	architecturalstyletypeid	90178	0.998926
5	basementsqft	90261	0.999845
6	bathroomcnt	58550	0.648574
7	bedroomcnt	58550	0.648574
8	buildingclasstypid	90267	0.999911
9	buildingqualitytypeid	70038	0.775829
10	calculatedbathnbr	58964	0.65316
11	decktypeid	90052	0.99753
12	finishedfloorlsquarefeet	87931	0.974035
13	calculatedfinishedsquarefeet	58778	0.651099
14	finishedsquarefeet12	60197	0.666818
15	finishedsquarefeet13	90261	0.999845
16	finishedsquarefeet15	89004	0.985921
17	finishedsquarefeet50	87931	0.974035
18	finishedsquarefeet6	90141	0.998516
19	fips	58550	0.648574
20	fireplacecnt	86924	0.96288
21	fullbathcnt	58964	0.65316
22	garagecarcnt	79830	0.884298
23	garagetotalsqft	79830	0.884298
24	hottuborspa	89479	0.991182
25	heatingorsystemtypeid	70512	0.78108
26	latitude	58550	0.648574
27	longitude	58550	0.648574
28	lotsizesquarefeet	62072	0.687588
29	poolcnt	84004	0.930534
30	poolsizeum	89944	0.996333
31	pooltypeid10	89887	0.995702
32	pooltypeid2	89867	0.99548
33	pooltypeid7	84412	0.935054
34	propertycountylandusecode	58550	0.648574
35	propertylandusetypeid	58550	0.648574
36	propertyzoningdesc	69685	0.771919
37	rawcensustractandblock	58550	0.648574
38	regionidcity	59216	0.655951
39	regionidcounty	58550	0.648574
40	regionidneighborhood	77632	0.85995
41	regionidzip	58562	0.648707
42	roomcnt	58550	0.648574

43	storytypeid	90261	0.999845
44	threequarterbathnbr	86021	0.952877
45	typeconstructiontypeid	90163	0.998759
46	unitcnt	69677	0.771831
47	yardbuildingsqft17	89364	0.989909
48	yardbuildingsqft26	90241	0.999623
49	yearbuilt	58810	0.651454
50	numberofstories	83076	0.920255
51	fireplaceflag	90181	0.998959
52	structuretaxvaluedollarcnt	58678	0.649992
53	taxvaluedollarcnt	58551	0.648585
54	assessmentyear	58550	0.648574
55	landtaxvaluedollarcnt	58551	0.648585
56	taxamount	58551	0.648585
57	taxdelinquencyflag	89662	0.99321
58	taxdelinquencyyear	89662	0.99321
59	censustractandblock	58758	0.650878

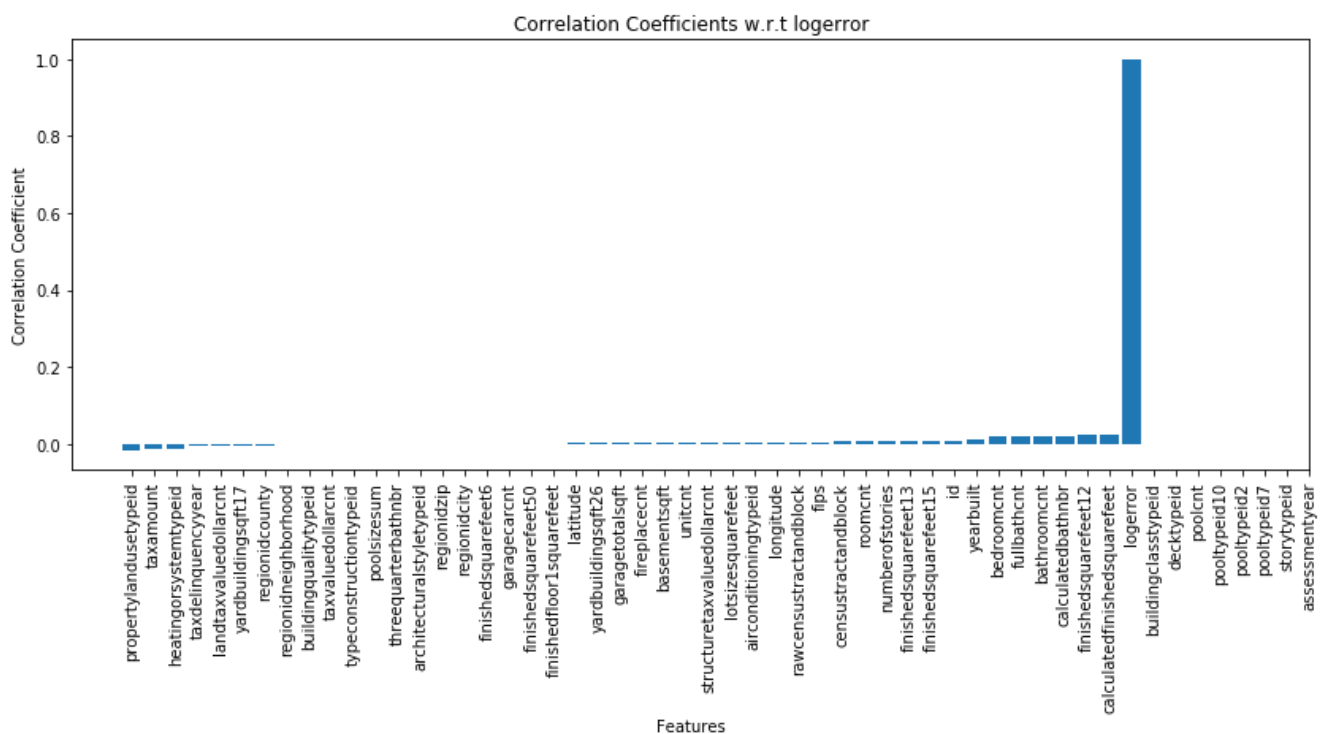
In [137]:

```
df_final=df_final.fillna(df_final.mean())
```

## c) univariate analysis

In [138]:

```
# make bar chart
correlation = df_final.corr().iloc[1,:]
correlation.sort_values(ascending=True, inplace=True)
plt.figure(figsize = (14,5))
plt.bar(np.arange(54),np.sort(correlation.values))
plt.title('Correlation Coefficients w.r.t logerror')
plt.xlabel('Features')
plt.ylabel('Correlation Coefficient')
plt.xticks(np.arange(54),correlation.keys(),rotation = 'vertical')
plt.show()
```



## explain reason

From the graph we can see that there are columns without any value of correlation. If we look at the data of these features, we find all the values of these features are the same. There is no change in values (variance) hence the correlation coefficient value with logerror is NaN resulting in missing correlation coefficient.

## d) non-linear regression model

In [139]:

```
df_final=df_final.drop(["hashottuborspa", "propertycountylandusecode", "propertyzoningdesc",
"fireplaceflag", "taxdelinquencyflag","id","transactiondate"], axis=1)
y=df_final.iloc[:,0]
x=df_final.iloc[:,1:]
```

In [140]:

```
# split and train
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(x,y, test_size=0.30)
```

In [141]:

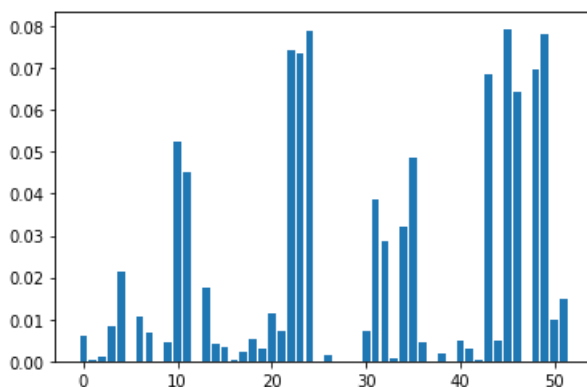
```
from sklearn.ensemble import RandomForestRegressor
rf = RandomForestRegressor()
rf.fit(X_train, y_train)
fi=rf.feature_importances_
y_pred=rf.predict(X_test)
```

C:\Users\mohan\Anaconda3\lib\site-packages\sklearn\ensemble\forest.py:245: FutureWarning: The default value of n\_estimators will change from 10 in version 0.20 to 100 in 0.22.  
"10 in version 0.20 to 100 in 0.22.", FutureWarning)

In [142]:

```
# report importances and mse
plt.bar(np.arange(len(fi)),fi)
from sklearn.metrics import mean_squared_error as mse
MSE1=mse(y_test,y_pred)
print(MSE1)
```

0.010475270588728618



## e) KFold

In [143]:

```
# KFold, k = 5
datax=np.array(x.iloc[:500,:])
datay=np.array(y.iloc[:500])
from sklearn.model_selection import KFold
kf=KFold(n_splits=5,shuffle=True)
MSE=[]
for a,b in kf.split(datax):
    trainx=datax[a,:]
    trainy=datay[a]
    testx=datax[b,:]
```

```
print('MSE for KFold:',MSE,'Avg MSE:',np.mean(MSE))
```

In [144]:

```
MSE_list
```

[illegible]

[illegible]

[illegible]



[illegible]



```

10 in version 0.20 to 100 in 0.22.", FutureWarning)
C:\Users\mohan\Anaconda3\lib\site-packages\sklearn\ensemble\forest.py:245: FutureWarning: The default value of n_estimators will change from 10 in version 0.20 to 100 in 0.22.
  "10 in version 0.20 to 100 in 0.22.", FutureWarning)
C:\Users\mohan\Anaconda3\lib\site-packages\sklearn\ensemble\forest.py:245: FutureWarning: The default value of n_estimators will change from 10 in version 0.20 to 100 in 0.22.
  "10 in version 0.20 to 100 in 0.22.", FutureWarning)
C:\Users\mohan\Anaconda3\lib\site-packages\sklearn\ensemble\forest.py:245: FutureWarning: The default value of n_estimators will change from 10 in version 0.20 to 100 in 0.22.
  "10 in version 0.20 to 100 in 0.22.", FutureWarning)
C:\Users\mohan\Anaconda3\lib\site-packages\sklearn\ensemble\forest.py:245: FutureWarning: The default value of n_estimators will change from 10 in version 0.20 to 100 in 0.22.
  "10 in version 0.20 to 100 in 0.22.", FutureWarning)
C:\Users\mohan\Anaconda3\lib\site-packages\sklearn\ensemble\forest.py:245: FutureWarning: The default value of n_estimators will change from 10 in version 0.20 to 100 in 0.22.
  "10 in version 0.20 to 100 in 0.22.", FutureWarning)
C:\Users\mohan\Anaconda3\lib\site-packages\sklearn\ensemble\forest.py:245: FutureWarning: The default value of n_estimators will change from 10 in version 0.20 to 100 in 0.22.
  "10 in version 0.20 to 100 in 0.22.", FutureWarning)
C:\Users\mohan\Anaconda3\lib\site-packages\sklearn\ensemble\forest.py:245: FutureWarning: The default value of n_estimators will change from 10 in version 0.20 to 100 in 0.22.
  "10 in version 0.20 to 100 in 0.22.", FutureWarning)
C:\Users\mohan\Anaconda3\lib\site-packages\sklearn\ensemble\forest.py:245: FutureWarning: The default value of n_estimators will change from 10 in version 0.20 to 100 in 0.22.
  "10 in version 0.20 to 100 in 0.22.", FutureWarning)
C:\Users\mohan\Anaconda3\lib\site-packages\sklearn\ensemble\forest.py:245: FutureWarning: The default value of n_estimators will change from 10 in version 0.20 to 100 in 0.22.
  "10 in version 0.20 to 100 in 0.22.", FutureWarning)
C:\Users\mohan\Anaconda3\lib\site-packages\sklearn\ensemble\forest.py:245: FutureWarning: The default value of n_estimators will change from 10 in version 0.20 to 100 in 0.22.
  "10 in version 0.20 to 100 in 0.22.", FutureWarning)

```

Out[144]:

```

[0.012061375227502174,
 0.01238404839507565,
 0.009757195579127425,
 0.009982592264602654,
 0.00971437781237707,
 0.013108918708190255,
 0.014059471380348455,
 0.014387240840575277,
 0.010612758897975855,
 0.009974540830847768,
 0.010939235820424357,
 0.010001761960604707,
 0.010056262720859542,
 0.006474531468530929,
 0.01251384120369906,
 0.007644469659575878,
 0.009421651911738157,
 0.009488837650338604,
 0.00942500872646759,
 0.008601529906074129,
 0.01364430968431742,
 0.008876130717341536,
 0.012079045482951984,
 0.011608889062973674,
 0.008581063169112045,
 0.0098001584848041,
 0.00852647815424317,
 0.008063283538660662,
 0.01230811828921441,
 0.010171141670555484,
 0.013186805373276805,
 0.009124226564229655,
 0.013089466763846657,
 0.007594882713935564,
 0.007493272174276747,
 0.014200948194069973,
 0.01173791198165334,
 0.0101991701688454,
 0.007482180298520437,
 0.012568526973940413,
 0.010071199285177936,
 0.010030386369471385,
 0.01173209160545139,
 0.008680805311553276,
 0.009242410377122886,
 0.012772586617029732,
 0.01349017594799377,
 0.015920203689387867,
 0.011356913324458272,
 ...

```

```
0.01066205190936414,  
0.013761994379965196,  
0.011481746891735234,  
0.007590576965291412,  
0.01057620419103451,  
0.010160726466442605,  
0.009434500685336128,  
0.012743664862470097,  
0.010108746780599785,  
0.009480276645307081,  
0.011953487365196304,  
0.007705424919251795,  
0.012476046939482488,  
0.013244709672437495,  
0.009804180750642825,  
0.008683195485729203,  
0.008076365347297231,  
0.011764994076166517,  
0.008843204335689733,  
0.011676443925342589,  
0.00845126794442038,  
0.012869709284217394,  
0.01038527085854723,  
0.013809789874356717,  
0.010916277233244074,  
0.0121205314820749,  
0.011567570528670667,  
0.009483813063117902,  
0.00790075924057199,  
0.011053234846193461,  
0.01092066588911933,  
0.011764254731947189,  
0.009870679277340753,  
0.01186040072747374,  
0.009476224367568347,  
0.009723240422728615,  
0.009266430904649937,  
0.010304594942030457,  
0.009144412915857714,  
0.01068377703270343,  
0.006882690887960816,  
0.011000535874831898,  
0.014355711716805341,  
0.014716527205212687,  
0.011581487273745705,  
0.009807680569712455,  
0.009745388416854388,  
0.010508018994221156,  
0.013541312634639103,  
0.013120929662676776,  
0.008980044725439814]
```

In [145]:

```
minMSE=min(MSE_list)  
maxMSE=max(MSE_list)  
print("Min MSE from list"+str(minMSE))  
print("Max MSE from list"+str(maxMSE))
```

```
Min MSE from list0.006474531468530929  
Max MSE from list0.015920203689387867
```

## Findings

MSE is different for each seed. This variation can affect the understanding of the results. To overcome this, Cross Validation is used. Cross-validation splits dataset into train and test in KFold and every data points get to be tested exactly once and is used in training k-1 times. This helps in generalization and reduces the bias.